

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
%matplotlib inline
```

```
boston_dataset=pd.read_csv("/content/BostonHousing.csv")
```

```
boston_dataset.head()
```

```

crim    zn    indus  chas    nox    rm    age    dis    rad    tax    ptratio    b    lstat    medv
0  0.00632  18.0    2.31    0  0.538  6.575  65.2  4.0900    1  296    15.3  396.90    4.98    24.0
1  0.02731    0.0    7.07    0  0.469  6.421  78.9  4.9671    2  242    17.8  396.90    9.14    21.6
2  0.02729    0.0    7.07    0  0.469  7.185  61.1  4.9671    2  242    17.8  392.83    4.03    34.7
3  0.03237    0.0    2.18    0  0.458  6.998  45.8  6.0622    3  222    18.7  394.63    2.94    33.4
4  0.06905    0.0    2.18    0  0.458  7.147  54.2  6.0622    3  222    18.7  396.90    5.33    36.2

```

```
boston_dataset.isnull().sum()
```

```

crim      0
zn        0
indus     0
chas      0
nox       0
rm        0
age       0
dis       0
rad       0
tax       0
ptratio   0
b         0
lstat     0
medv      0
dtype: int64

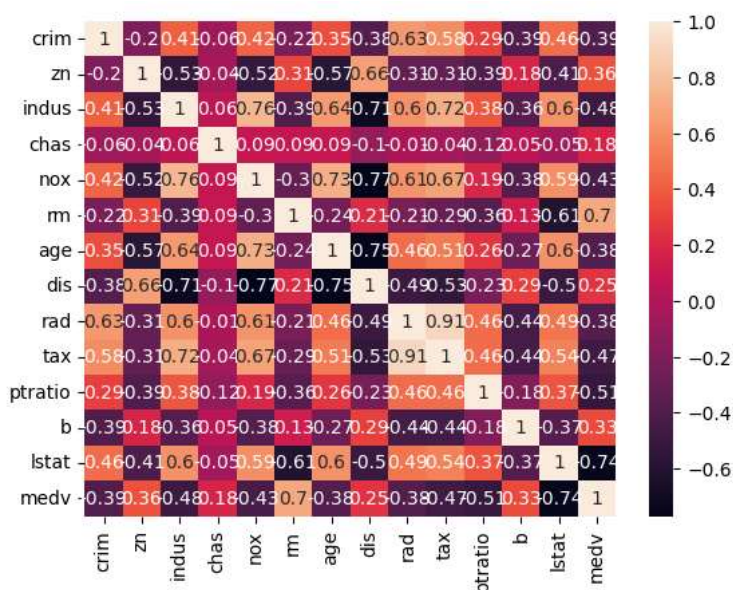
```

```

correlation_matrix = boston_dataset.corr().round(2)
sns.heatmap(data=correlation_matrix,annot=True)

```

<Axes: >



```
plt.figure(figsize=(20, 5))
```

```



features = ['lstat', 'rm']
target = boston_dataset['medv']

```

```

for i, col in enumerate(features):
    plt.subplot(1, len(features) , i+1)
    x = boston_dataset[col]
    y = target
    plt.scatter(x, y, marker='o')
    plt.title(col)
    plt.xlabel(col)
    plt.ylabel('medv')

```



```

X = pd.DataFrame(np.c_[boston_dataset['lstat'], boston_dataset['rm']], columns = ['lstat','rm'])
Y = boston_dataset['medv']

```

```

from sklearn.model_selection import train_test_split

```

```

X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size = 0.2, random_state=5)
print(X_train.shape)
print(X_test.shape)
print(Y_train.shape)
print(Y_test.shape)

```

```

(404, 2)
(102, 2)
(404,)
(102,)

```

```

from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error
import statsmodels.api as sm

```

```

lin_model = LinearRegression()
lin_model.fit(X_train, Y_train)

```

```

▼ LinearRegression
LinearRegression()

```

```

model = sm.OLS(Y_train, X_train).fit()

```

```

print(model.summary())

```

```

=====
                        OLS Regression Results
=====
Dep. Variable:          medv    R-squared (uncentered):          0.947
Model:                  OLS    Adj. R-squared (uncentered):      0.947
Method:                 Least Squares    F-statistic:          3581.
Date:                   Thu, 29 Jun 2023    Prob (F-statistic):    6.67e-257
Time:                   14:36:10    Log-Likelihood:        -1272.2
No. Observations:       404    AIC:                   2548.
Df Residuals:           402    BIC:                   2556.

```

```
Df Model:                2
Covariance Type:         nonrobust
=====
              coef      std err          t      P>|t|      [0.025      0.975]
-----
lstat      -0.6911      0.036     -19.367      0.000     -0.761     -0.621
rm         4.9699      0.081      61.521      0.000       4.811       5.129
=====
Omnibus:            121.894   Durbin-Watson:           2.063
Prob(Omnibus):      0.000   Jarque-Bera (JB):       389.671
Skew:              1.370   Prob(JB):           2.42e-85
Kurtosis:          6.954   Cond. No.           4.70
=====
```

Notes:
[1] R² is computed without centering (uncentered) since the model does not contain a constant.
[2] Standard Errors assume that the covariance matrix of the errors is correctly specified.